



(a) Full MaskedBiMamba: 1 s window, T = 30 samples



x and y: per-frame min-max normalization; c: joint confidence. Batch dimension omitted.

(b) Independent forward and backward Mamba stacks

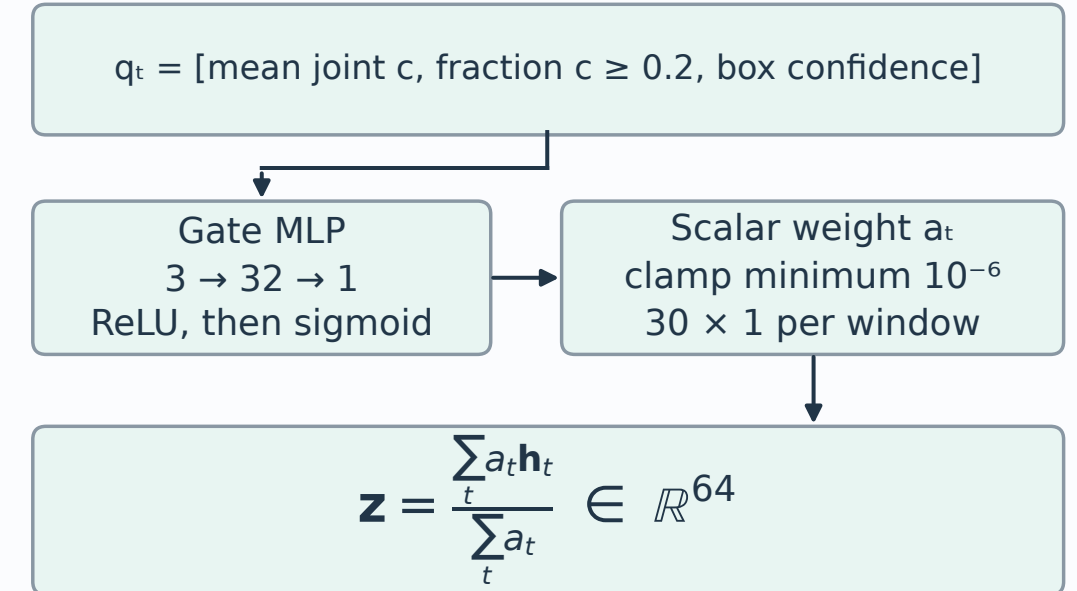
Both stacks: 2 residual layers; d_model = 64, d_state = 16, d_conv = 4, expand = 2



Each residual layer: $U^{(l+1)} = U^{(l)} + \text{Mamba}_l(U^{(l)})$

Reverse outputs are realigned before fusion; both directions use the complete window.

(c) Quality gate and temporal pooling



q_t bypasses the training-time pose mask.

Masking zeroes the full 36-D pose vector independently per frame; at least one frame is retained. It is disabled for evaluation.